How Far Can a Digital Twin of a Low-Cost Robot Be Trusted? A Preregistered Comparison of a Coupled Electrical and Dynamic Twin with Two Real Robots

Muthuphei Mukhunyeledzi, Branden Ingram and Pravesh Ranchod
School of Computer Science and Applied Mathematics
University of the Witwatersrand, Johannesburg, South Africa

***Abstract*—A robot moves only as well as its circuit lets it, yet robot simulators treat a motor as an ideal torque source and digital twins of small robots are rarely tested against the hardware they copy. We build a digital twin of the Waveshare JetBot in which a circuit-level model of the power electronics, written in *elecsim*, a code-first simulation framework we develop, is solved together with the robot's dynamics at every step. We verify elecsim against ngspice and test the twin against two physical robots with predictions saved in advance and a preregistered fifty-run test. Our research shows that the coupled twin predicts steady motion to within a few percent, including held-out spins on a second robot to within 0.5%. An ablation shows what the coupling contributes: a twin with an ideal supply, calibrated the same way, misses the same spins by 4.9% because it cannot follow the battery's charge, and cannot predict the Jetson's supply falling to 4.53 V, below its specification, which the coupled twin predicts to 4 mV. Over fifty straights the twin passes the preregistered test yet is 8% short on every run that goes straight; separating the caster's rolling and scrubbing friction reduces this to 0.9%, and cross-validation over all batches keeps it within 1.3%. Short open-loop turns are not predicted on either robot, and the robots' occasional swerves are traced to a sustained wheel-speed difference, not a timing fault.**

***Index Terms*—digital twin, reality gap, electromechanical co-simulation, model validation, differential-drive robot, preregistration**

I. Introduction

A *digital twin* is a virtual model of a specific physical system that is kept close enough to it to stand in for it: to predict its behaviour, to test changes before they are made, or to detect when the real system departs from what is expected [Grieves and Vickers 2017; Tao *et al.* 2019]. Reviews of the field find that few published twins exchange data with their physical system in both directions [Kritzinger *et al.* 2018], and that the benefits claimed for twins are rarely validated or quantified [Jones *et al.* 2020]. In robotics the mismatch between simulated and real behaviour is the *reality gap* [Jakobi *et al.* 1995; Koos *et al.* 2013]. It is usually worked around by randomising the simulator [Tobin *et al.* 2017; Peng *et al.* 2018], and there is no common benchmark for measuring it [Muratore *et al.* 2022].

A physical robot is an electrical system as much as a mechanical one. Its motors turn only as fast as the battery, the protection circuit, the switch and the driver chip allow; when the battery runs down or the processor draws more power, the motion changes. Robot simulators nevertheless drive a motor by a force, torque or velocity command [Koenig and Howard 2004; Todorov *et al.* 2012], without its voltage, current and back-EMF, and circuit simulators stop at the motor terminals [Nagel and Pederson 1973]. We build a twin in which the two are one system: at every step the circuit is solved for the motor currents given the wheel speeds, and the dynamics move the robot given those currents. The circuit model shows whether the electronics are working, and the coupling shows what their state does to the robot's motion.

We study the question on a small, inexpensive robot: *how far can a digital twin of a low-cost mobile robot be trusted, where does it fail, and what does modelling the circuit add?* The platform is the Waveshare JetBot AI Kit [Waveshare 2019], a two-wheel educational robot built around an NVIDIA Jetson Nano [NVIDIA AI IOT 2019]. The project began as the design of a RoboCup Small Size League robot [Kitano *et al.* 1997]; when parts could not be obtained in time, it was redirected to a robot that exists, so that the twin can be judged by whether it is right rather than by whether it looks plausible. Our contributions are:

1) *elecsim*, a Python framework in which a circuit, its ratings and its requirements are code, verified against ngspice (Section III).
2) A JetBot twin that couples the board's power path to the robot's dynamics and exposes the robot's own software interface.
3) A validation protocol that keeps the comparison a test rather than a fit: provenance, saved predictions, staged calibration, held-out and preregistered tests, and cross-validation (Section IV).
4) Results on two robots that locate where the twin can be trusted, an ablation that isolates what the electrical

coupling contributes, and two model diagnoses, of the caster and of the swerves (Section V).

II. Related Work

A. Digital twins

Grieves and Vickers [2017] define a digital twin as a virtual representation linked to a physical product throughout its life. Kritzinger *et al.* [2018] distinguish a *digital model*, a *digital shadow* (data flow from the physical object only) and a full twin (data flow both ways), and find that most manufacturing "twins" are models or shadows. Later reviews name levels of fidelity [Jones *et al.* 2020] and model validation [Fuller *et al.* 2020] as open problems, and Rasheed *et al.* [2020] argue that physics-based models generalise to problems governed by the same physics but can carry large errors from uncertain inputs, so they should be combined with measured data. In these terms our twin is a digital shadow: the robot's logs condition every replay, but the twin does not command the robot. We focus on measured agreement between the twin and its robot.

B. The reality gap

Jakobi *et al.* [1995] popularise the term and show that realistic noise in a simulation helps evolved controllers transfer to the real robot; Koos *et al.* [2013] learn which simulated behaviours transfer. Sim-to-real work in learning-based control randomises the simulator's appearance [Tobin *et al.* 2017] or dynamics [Peng *et al.* 2018; Muratore *et al.* 2022] so that what is learned is robust to the gap, or improves the actuator model, analytically with a DC-motor model [Tan *et al.* 2018] or with a network learned from data [Hwangbo *et al.* 2019]. Surveys [Zhao *et al.* 2020; Muratore *et al.* 2022] catalogue the methods for crossing the gap, chiefly domain randomisation; Muratore *et al.* [2022] note that there is no common benchmark for sim-to-real transfer and that physics-engine comparisons cover few scenarios, and Choi *et al.* [2021] identify validated simulation, and models of contact and friction, as central open challenges. We do not train a controller; we measure the gap of one deterministic model term by term.

C. Simulators and circuit models

Gazebo [Koenig and Howard 2004] and MuJoCo [Todorov *et al.* 2012] model rigid-body dynamics and contact; no engine is best on every test [Erez *et al.* 2015], and contact remains hard to simulate accurately [Erez *et al.* 2015; Collins *et al.* 2021]. None models the electrical supply. Circuit simulators descend from SPICE [Nagel and Pederson 1973] and solve modified nodal analysis (MNA) [Ho *et al.* 1975]. Acausal languages such as Modelica [Fritzson 2014] can couple electrical and mechanical domains but need a dedicated toolchain. elecsim uses MNA, but parts, ratings and requirements are ordinary Python, so a design is tested and versioned like any other code and can be stepped inside a robot simulation loop.

D. Calibration and validation

System identification [Ljung 1999] estimates model parameters from experiments. Kennedy and O'Hagan [2001] show that even the best-fitting parameters of an inadequate model leave a systematic discrepancy, and Brynjarsdóttir and O'Hagan [2014] show that calibrating without accounting for it biases the parameters: a wrong model compensates by moving the wrong numbers. The verification and validation literature [Oberkampf and Roy 2010; Sargent 2013] separates calibration from validation data. For wheeled robots, Borenstein and Feng [1996] separate systematic from non-systematic odometry errors by repeated runs, and Siegwart *et al.* [2011] list wheel slip and uneven floor contact among the non-systematic error sources. Preregistering hypotheses and analyses [Nosek *et al.* 2018] guards against the data-dependent choices that make an analysis agree with its data [Gelman and Loken 2014]. We adopt it for the main test.

III. System

A. elecsim: a code-first electrical framework

A design in elecsim is a Python module whose `build()` function returns a circuit. A *component* is a class that declares its ports and parameters and supplies `stamp`, which adds its linearised equations, and `measure`, which names what can be observed about it. A *circuit* joins components into nets. *Constraints* are of three kinds: *rules* check the wiring without simulating, *limits* compare each signal with a component's rating at every step, and *requirements* ask whether a finished run met its purpose ("the 5 V rail stays above 4.75 V"). A *report* lists every constraint in every scenario with the unverified numbers it rests on: every parameter carries its provenance, `datasheet`, `measured` or `assumed`.

The solver assembles and solves, at each Newton iteration, $\mathbf{G}(\mathbf{x}_k)\,\Delta\mathbf{x} = -\mathbf{f}(\mathbf{x}_k)$ [Ho *et al.* 1975], with $\mathbf{x}$ the node voltages and any branch currents, and steps in time by backward Euler with automatic step halving. A damped retry and a fallback in which regulators that lose their input switch off make it robust to parts written by non-specialists. The library covers passives, sources, diodes, MOSFET switches, a lithium cell with a state-of-charge curve, a pack protector, a buck converter, a DC motor, an H-bridge, the PCA9685 PWM controller and measurement chips with their real resolution. An averaged H-bridge solves the periodic winding current in closed form, correct in discontinuous conduction and in both decay modes, and within 3% of the switching model; with it the twin runs at about three times real time. Sixty automated tests cover the framework.

B. The coupled JetBot twin

The expansion board's schematic is transcribed component by component: three 18650 cells and their holder, the protection circuit, the current-sense shunt and INA219 monitor, the slide switch, the buck converter feeding the Jetson, the PWM controller and the motor drivers. The board has footprints for two alternative driver chips, a TB6612FNG or two DRV8870s, which treat the motor differently between PWM pulses; the

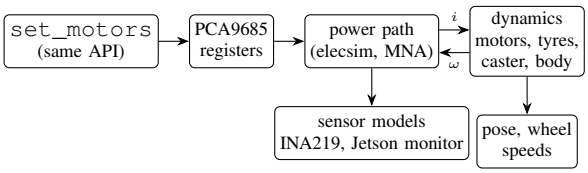


Fig. 1. The coupled twin. A motor command becomes the same PWM register writes the real library makes. Each 1 ms step the circuit is solved for the motor currents i given the wheel speeds ω (back-EMF), and the dynamics advance the robot under the resulting torques.

twin contains both. The Jetson Nano is written as a new component, the intended way to extend the framework.

The twin's `Robot` object has the robot library's methods: `set_motors(left, right)` writes the same PCA9685 registers as the real `Motor` class, and a test asserts identical register values. The circuit and the dynamics are coupled in both directions (Fig. 1): the wheel speeds set the motors' back-EMF, the circuit returns the winding currents, and these set the torques. A run-down battery, a heavier Jetson load or a different driver chip therefore changes the motion through the same equations as on the robot.

The dynamic state is position, heading, forward speed v, yaw rate ω_z and the two wheel speeds. Each motor follows the DC model [Hughes and Drury 2019], reflected through the gearbox with an efficiency, and each wheel carries the rotor inertia and Coulomb and viscous friction. Each tyre transmits a force that saturates at its friction limit, $F = \mu_t N \tanh((r\omega - v_w)/s_0)$, a smooth approximation of the saturating slip curve [Pacejka 2012]. The caster is a ball at distance x_c behind the axle whose contact point moves at $(v,\, \omega_z x_c)$ in the body frame. As first built, it resists with one coefficient μ_c in every direction; Section V-E shows why this is replaced by

$$f_x = -\mu_{\text{roll}} N_c \tanh\frac{v}{\epsilon}, \qquad f_y = -\mu_{\text{scrub}} N_c \tanh\frac{\omega_z x_c}{\epsilon}, \quad (1)$$

with separate resistance to rolling along the heading and to scrubbing across it; with μ_{roll} unset the model reduces to the original. Load transfer between wheels and caster is included; sideways tyre slip, gearbox backlash and differences between the motors are not. Geometry comes from the manufacturer's CAD model, which also drives a browser viewer (Fig. 2) that shows the robot moving beside the power path's live currents and voltages and every part against its rating, so the effect of the circuit on the motion can be watched as it happens.

IV. Methodology

A. Research questions

- RQ1 How closely does the twin, built from documentation and calibrated in isolated stages, predict the robot before each run?
- RQ2 Does the twin fall within the robot's own variation over many repeats of one manoeuvre?
- RQ3 Where the twin fails, can the cause be located in the model?
- RQ4 What does coupling the circuit to the dynamics add?

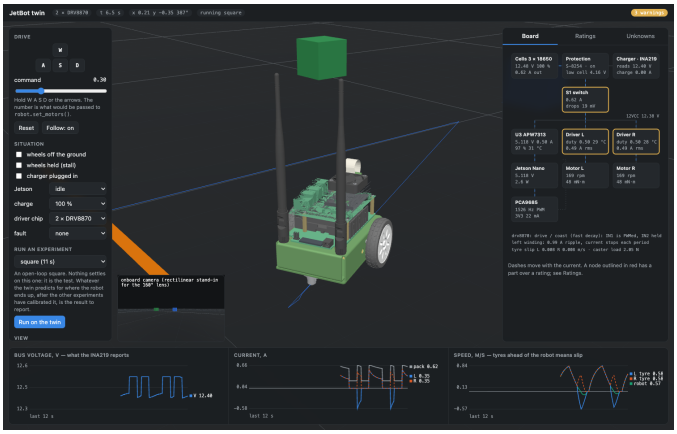


Fig. 2. The live twin in the browser: the manufacturer's CAD, the board's power path with live currents and voltages, and traces of what the robot's own sensors would report. All values come from the coupled Python twin.

B. What can be observed

A twin can only be compared with a robot through the robot's sensors. The JetBot measures its battery rail voltage (INA219) and the Jetson reports its own input voltage and power. The INA219's shunt lies between the charger jack and the battery rail, so the robot cannot measure its own load or motor current, wheel speed or position. The twin reports what each real sensor would report, at its resolution; everything else is measured from video: wheel speed from a marker on the hub, paths using the 50 cm floor tiles as a ruler, and turn angles from the robot's boards at rest.

C. Protocol

Provenance. Every number is tagged as datasheet, measured or assumed, and the assumed ones are printed with every result. *Predictions before measurement.* Before the robot is powered, forty twins drawn from the uncertainty of the assumed parameters give a 5–95% band for every experiment; before each floor run, the calibrated twin's prediction is saved. *Staged calibration.* Each experiment isolates a few parameters (Table I), fitted by least squares with the rest held. Geometry is measured and never fitted, and the driver chip is identified by a direct electrical test, because a rehearsal shows that a free fit chooses both wrongly. *Held-out tests.* An open-loop square fits nothing; on the second robot, the fifty straights fit nothing in the first round. Each real run is compared with the twin *replayed* on that run's conditions: the command times as executed, the Jetson power logged, and the state of charge implied by the resting voltage.

Rehearsal. Before touching hardware, the procedure is run against a second twin with nine parameters moved by 10–45%, seen through noisy sensors. It exposes two faults: with the wheel radius free, the fit returns a 10 mm wheel while still scoring well on the held-out square [Brynjarsdóttir and O'Hagan 2014], and an inferred driver chip is the wrong one. After both fixes the procedure recovers every altered value to within 9%.

TABLE I
EXPERIMENTS AND THE PARAMETERS EACH ISOLATES

Experiment	Set-up	Settles
CPU load step	bench	pack and 5 V path resistance
driver test	wheels up	which driver chip is fitted
wheels-up sweep	wheels up	motor constant, friction
stall	wheels held	winding resistance
spin	floor	caster friction
straight	floor	drag on the floor
square, short turns	floor	nothing (held out)
50 straights (robot 2)	floor	nothing (preregistered)

D. The preregistered test on robot 2

After robot 1 develops a fault, a second unit is calibrated from its own bench and spin experiments. Before its first straight run we commit the prediction and the test: both wheels at 0.5 for 1.5 s from rest, fifty times in ten batches of five, filmed from 1.01 m. The planned analyses are (1) a spread test, passed if zero lies within the middle 95% of the residuals $d_i = \text{robot}_i - \text{twin}_i$; (2) the bias, with a 95% t-interval over batch means; (3) spread within and between batches; and (4) drift and heading, reported only. No run is excluded. Positions come from a video tool that uses the tile joints as a ruler (2.4 mm RMS) and corrects for the robot's height using its LED, 8.5 cm above the floor, leaving about $\pm 0.5\%$ uncertainty.

E. Coupling ablation and cross-validation

To answer RQ4 we build three twins that differ only in what the motors are fed: the *coupled* twin; a *no-sag* twin, whose power path has no resistance but still follows the battery's charge; and an *ideal-supply* twin, also fixed at full charge, as a robot simulator without an electrical model would assume. Each goes through the same calibration (motor constant from the wheels-up sweep, caster scrubbing from two spins, rolling from straights 1–5) and the same held-out tests, so none starts at a disadvantage. To test whether the caster fit depends on which runs it uses, we repeat it with each batch of five as the fit set and the other 45 runs as the test.

V. RESULTS

A. Verifying elecsim against SPICE

Table II compares elecsim with ngspice on five circuits, the last a copy of the JetBot's power path with robot 2's measured values as both stalled motors are switched on. DC solutions agree to solver precision. In transients the error falls in proportion to the step, as backward Euler should (slope 0.98 on the RC circuit), and the method's known numerical damping shows on the lightly damped RLC circuit. At the twin's 1 ms step, elecsim misses the winding's first millisecond of current rise, whose time constant is 0.6 ms; within 5 ms it agrees with SPICE to 4 mV on a 2.5 V swing. Reducing the twin's step eightfold changes its predicted straight by 0.2 mm and its spin by 0.2°, so the motion predictions do not depend on the step.

TABLE II
ELECSIM AGAINST NGSPICE

Circuit	Step	Largest difference
Resistor bridge (DC)	–	10^{-8} V
Diode and resistor, 0–5 V	–	0.04% of current
RC charge	$\tau/100$	0.2% of swing
RLC ring-down	$T/80$	3.4% of swing
JetBot power path	1 ms	4 mV after 5 ms

B. Electrical behaviour (RQ1, RQ3)

The driver chip is identified without opening the robot. Driving the DRV8870 channels moves nothing; driving the TB6612FNG channels dips the rail by 204 mV. The board carries the TB6612FNG, contrary to our initial assumption, and, as the twin predicts for that chip, the lowest command that turns a free wheel is 0.10.

The Jetson's supply sags below its specification. Under a CPU load step, which the Jetson's own power monitor makes a known load, the twin built from the schematic holds the Jetson at 5.12 V while the Jetson reports 4.50 V, below its 4.75 V minimum. Fitting a 0.49 Ω resistance between the regulator and the Jetson, through header pins and carrier traces on no schematic, reduces the RMS rail error from 80 to 30 mV. Robot 2 sags to 4.53 V through 0.52 Ω, so the sag is a property of the design. The reading is the Jetson's own and is not yet confirmed with a meter. The pack's resistance also doubles as it empties (0.058 to 0.125 Ω per cell), which the twin treats as a constant.

C. Motion on robot 1 (RQ1)

With the wheels in the air, the twin as designed is 14% from the measured wheel speed and 8% after calibration; the predicted bands hold (199–284 rpm at a command of 0.5, measured 268). Spin-up and braking times, never fitted, agree: 80 against 83 ms and 140 against 162 ms. Above a command of about 0.8 (0.6 on robot 2) the robot's speed stops rising in proportion, which the motor model cannot follow, so robot 2's motors are fitted on 0.1–0.6. On the floor, the twin's saved predictions are 1.47 m for a 2 s straight (measured 1.43 m, 3%) and 1095° for a 2 s spin (1033°, 6%).

On the held-out square the twin gets every side to within 3 cm and misses the corners by up to 39°. Repeating the 0.29 s corner from rest shows why in part: a motor command is fourteen I^{2}C writes from Python, so its duration varies (254 to 362 ms) and the angle follows it (correlation 0.82). Fixing the agent's timing reduces the variation from 29 to 2.2 ms. Twenty-four timed turns then average 94.3° (SD 7.6) against the twin's 83.7° (Fig. 3a), a 12.6% bias that no single parameter explains without breaking the spin or the straight; corners taken straight after a side vary three times as much (SD 23°). A strip of tape left on one tyre is found late; removing it changes distance and rotation by under 3% but reveals a 2.6% mismatch between the motors that the twin, with identical motors, cannot produce.

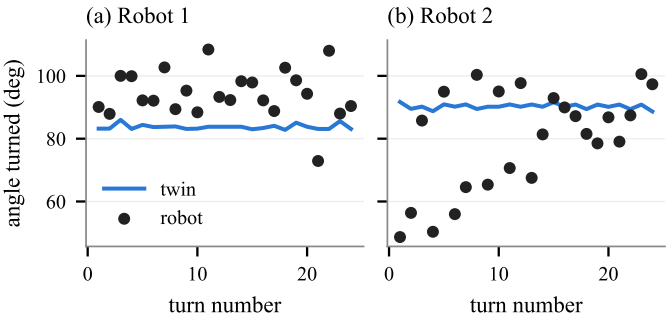


Fig. 3. Twenty-four identical short turns (0.29 s at ∓0.3, from rest), robot against the twin replayed on each turn's own timing. Robot 1 turns 12.6% further than its twin; robot 2 turns 11.6% less and scatters twice as much.

TABLE III
FIFTY STRAIGHTS ON ROBOT 2 (ROBOT MINUS TWIN)

	Round 1 (50 runs)	Round 2 (runs 6–50)
Spread test	pass	pass
Mean bias	+6.0%	−1.5%
95% interval (cm)	+4.6 to +8.0	−3.5 to +0.2
Runs that go straight	+8.2% (40)	+0.9% (35)
Runs that swerve (10)	−2.9%	−9.6%

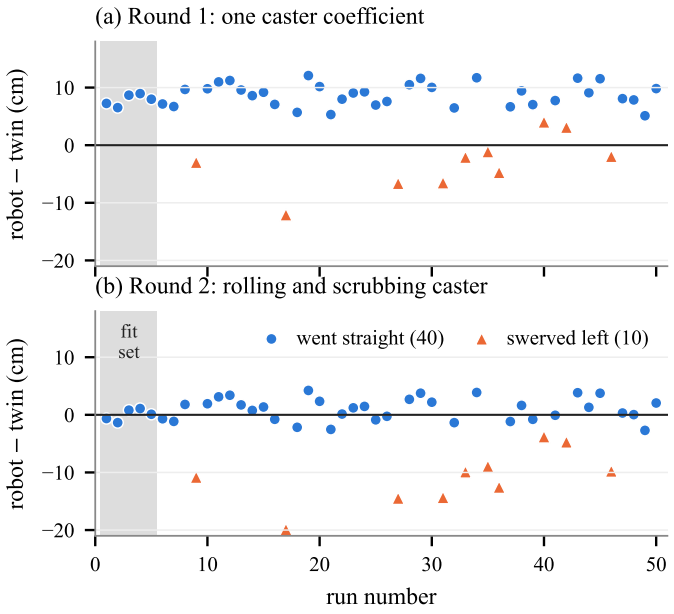


Fig. 4. Robot minus twin for robot 2's fifty straights. (a) The twin as calibrated on bench tests and spins. (b) The twin with the caster split into rolling and scrubbing friction, fitted on runs 1–5 only (shaded). Triangles mark runs that swerve more than 25 cm left.

D. Fifty straights on robot 2 (RQ2)

Robot 2 differs measurably from robot 1: 29% higher cell resistance, 9% lower winding resistance, and a caster that drags half again as much (0.506 against 0.324). Table III gives the preregistered results. The robot travels 1.120 m on average and the twin 1.056 m; zero lies within the middle 95% of the residuals (−6.6 to +11.7 cm), so *the twin passes the preregistered spread test*, with a bias of +6.0% (interval +4.6 to +8.0 cm). Batches do not differ beyond chance.

The data show what the plan does not anticipate (Fig. 4a). In 10 of the 50 runs the robot swerves hard left on open floor, ending 26 to 49 cm to the side. We report this split as exploratory. On the 40 runs that go straight, the robot repeats itself to 1.8 cm and *the twin is 8.2% short on every one*; on these alone the spread test would fail. The test passes because the swerves widen the robot's spread.

E. Locating and fixing the error (RQ3)

Robot 2's caster coefficient is fitted on spins, where the ball is dragged sideways; driving straight, it rolls. Replaying one run with different caster coefficients gives 1.061 m at 0.506 (the spin fit), 1.113 m at 0.324 (robot 1's) and 1.198 m at 0.02 (a free ball), against 1.144 m measured, while a 10% change in gearbox efficiency moves it by only 7 mm. On robot 1 the mismatch is hidden because gearbox efficiency and caster friction are fitted together on a straight and a spin, and the efficiency moves further (0.65 to 0.51) than a gearbox plausibly would to absorb it, as Brynjarsdóttir and O'Hagan [2014] warn.

We therefore separate rolling and scrubbing (Eq. 1), committing the fit and test before refitting: μ_{scrub} from two spins (0.524), μ_{roll} from runs 1–5 (0.227). On the 35 held-out runs that go straight, the error falls from 8.2% to 0.9% (Fig. 4b) and the spread test passes on them alone. Cross-validation shows that this does not depend on the choice of runs 1–5: fitted on the straight runs of each batch in turn, μ_{roll} lies between 0.15 and 0.23 and the error on the other batches' straight runs between −1.3 and +1.0% (nine folds; one batch has too few straight runs). A 95% prediction interval from each fit set's own scatter covers 97% of held-out straight runs on average, but only 80% of all runs, because the swerves fall outside it. Fitted on batches that include swerves, μ_{roll} rises to as much as 0.45. The fitted value is above a free ball's and absorbs losses the twin does not separate, such as tyre rolling resistance.

Ten new spins with the corrected twin average 1039° against the twin's 1044° (−0.5%, interval −16 to +6°). Twenty-four short turns average 79.8° (SD 16.2) against 90.3°, a −11.6% bias (Fig. 3b): robot 2 turns less than its twin where robot 1 turns more, so the short-turn error is not one systematic term.

F. What the electrical coupling adds (RQ4)

Table IV gives the ablation. All three twins fit the wheels-up sweep equally well (1.1% RMS), and removing the power path's resistance alone changes no motion prediction: at these currents the sag is too small to matter for motion. Following the battery's charge does matter. The sweep is recorded at 11.5 V, sagging to 10.2 V under load, and the test runs at 12.2–12.4 V. The ideal-supply twin, which cannot tell these apart, learns a motor 10% too weak; on the ten held-out spins it is 4.9% short, against 0.5% for the coupled twin, and on the straights it remains 2.3% short even with its rolling friction driven to zero, so calibration cannot repair it. Only the coupled

TABLE IV
ABLATION ON ROBOT 2: WHAT THE MOTORS ARE FED (ROBOT MINUS TWIN)

	Coupled	No sag	Ideal supply
Motor speed at 6 V (fitted, rpm)	249	248	225
Caster rolling (fitted)	0.230	0.252	0 (bound)
35 held-out straights	+0.9%	+0.9%	+2.3%
10 held-out spins	−0.5%	−0.6%	+4.9%
24 short turns	−11.4%	−11.2%	−9.5%
Bus sag, straights (robot 149 mV)	81 mV	0	0
Jetson 5 V loaded (robot 4.53 V)	4.53 V	5.12 V	5.12 V

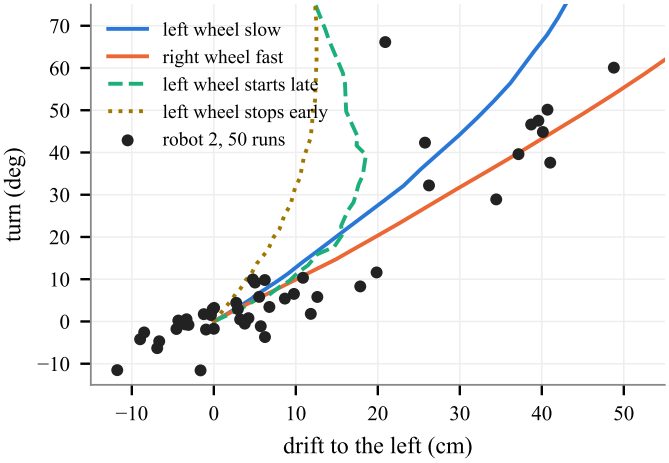


Fig. 5. Drift against turn for robot 2's fifty straights, with the path of each swerve hypothesis traced in the twin. Only a sustained difference between the wheels passes through the ten swerves. Two turn readings that are known misreads (92° and 147°) lie above the plot.

twin predicts the electrical state: the Jetson's supply under load to 4 mV, and a bus sag during each straight, although it predicts 81 mV where the robot shows 149 mV, which suggests the twin underestimates the current drawn on the floor.

G. Why the robot swerves (RQ3)

The shape of a swerve shows when the turning happens. If one wheel is slower for the whole run, the path is an arc and the sideways drift is about half of distance times turn; if a wheel starts late, the robot turns first and the drift is about the full product; if a wheel stops early, it turns last and drifts little. We trace each hypothesis in the twin and match every swerve to the nearest point on each curve, in units of the video's measurement error (5 cm, 5°), measuring the twin's path exactly as the video tool measures the robot's (Fig. 5). A sustained speed difference of about 30% between the wheels matches 9 of the 10 swerves within two measurement errors (RMS 0.9); a late start matches 2 and an early stop none (RMS 4.2 and 5.3). The swerves' distance, 1.02 m, lies between that implied by a slow left wheel (0.84 m) and a fast right one (1.19 m), so the data do not say which wheel. The cause is therefore something that persists through the run, such as traction or the caster, not command timing.

TABLE V
TWIN AGAINST ROBOT: SUMMARY (ROBOT MINUS TWIN)

Test	Result	Before run?
Wheel speed, wheels up	8%	bands held
Spin-up / braking time	80/83, 140/162 ms	never fitted
R1 straight / 2 s spin	3% / 6%	yes
R1 square: sides / corners	3 cm / 39°	yes
R1 / R2 short turns	+12.6% / −11.6%	yes
R2 50 straights, planned test	pass; +6.0%	preregistered
R2 straights, caster fixed	0.9% (35 held out)	preregistered fit
cross-validated	−1.3 to +1.0%	9 folds
R2 10 spins	−0.5%	preregistered
Jetson 5 V under load	4 mV	fitted
Swerves	30% wheel difference	diagnosed

VI. DISCUSSION

Where the twin can be trusted (RQ1). Our research shows that a physically based twin of a low-cost robot, built from its documentation and calibrated in isolated stages, predicts steady behaviour to within a few percent (Table V): electrical margins, wheel speed, transient timing it is never fitted to, spin angle, and straight-line distance once the caster has the right structure. These are what a designer asks of a twin: whether a supply is adequate, how fast the robot moves, what a different driver chip or a flat pack would change.

Why the circuit belongs in the twin (RQ4). The ablation shows that the coupling earns its place for two reasons. The motion depends on the battery's state, and a twin that cannot follow it is miscalibrated by an amount no friction coefficient can absorb. And the electrical state is itself a prediction a robot simulator cannot make: the coupled twin finds the Jetson running below its rated voltage under load, which no motion test would reveal. The ablation also shows the limits of the claim: on a charged pack at moderate commands, the supply's internal resistance does not change the motion, so the gain comes from the battery's charge, not its sag.

Where it cannot be trusted. Short open-loop turns are not predicted on either robot, and the robots do not repeat them: robot 1 scatters by ±8° and robot 2 by ±16°, with biases of opposite sign. A 0.3 s turn is almost all acceleration and braking, where backlash, the caster's swivel and bus timing dominate, none of which a lumped model of this kind captures; the honest output for such manoeuvres is a distribution. The swerves remain beyond a deterministic twin, but they are now characterised: a sustained wheel-speed difference of about 30%, appearing in one run in five.

A pass is not the whole answer (RQ2). The preregistered test is passed, and legitimately, but only because swerves the twin cannot produce widen the robot's spread; on the runs the twin does describe, it is 8% short every time. Reporting the bias and the exploratory split alongside the planned test is what locates the error. The cause lies in the model's form, one friction for two kinds of motion, which calibration alone cannot repair [Kennedy and O'Hagan 2001] and which randomising a single friction coefficient [Peng *et al.* 2018] would not cover; this

matches the view that contact and friction remain among the hardest phenomena to simulate [Erez *et al.* 2015; Choi *et al.* 2021].

Threats to validity. We test two robots on one floor, and robot 2's fifty runs are made in one 23-minute session. Most bench quantities are measured once or twice. The swerve split and the form of the caster correction are chosen after seeing the data, though cross-validation shows the fitted value is stable. The swerve diagnosis assumes that the video tool's drift and turn are accurate to 5 cm and 5°. Motor current is never measured directly, and the twin's under-predicted bus sag suggests the floor current is underestimated. The robot's mass is estimated from its parts list, and robot 1's 24 turns are made with tape on one tyre. Several of our own errors are caught during the work, among them the wheel-radius fit, a logging fault in the robot's agent, and a height correction that understates every distance by 4% until it is checked against the LED; each is caught by a check built for the purpose.

VII. Conclusion

We build a digital twin of the Waveshare JetBot in which a circuit model, written in the code-first framework elecsim and verified against ngspice, is solved together with the robot's dynamics, and we test it against two robots with saved predictions and a preregistered fifty-run test. Our research shows that the twin can be trusted for steady motion and the electrical budget, to within a few percent, and cannot be trusted for short open-loop turns, which the robots themselves do not repeat. Modelling the circuit is what lets the twin follow the battery's charge, a 4.9% difference on held-out spins, and predict the Jetson's under-voltage, which the schematic hides.

The main methodological lesson is that a twin which passes its planned test can still be consistently wrong, and the error is found only by repeating the experiment on a second unit and looking past the verdict to the bias. Once found, it lies in the form of the model, and fixing the form, not retuning a number, removes it on held-out and cross-validated runs.

Future work should measure motor current directly to resolve the under-predicted bus sag; model pack resistance as a function of charge; separate backlash from post-braking state as causes of the turn error; identify which wheel slows in a swerve with a marker on each hub; and repeat the fifty straights on other days and on robot 1. The framework, twin, data and analysis scripts are in the project repository.

References

[Borenstein and Feng 1996] Johann Borenstein and Liqiang Feng. Measurement and correction of systematic odometry errors in mobile robots. *IEEE Transactions on Robotics and Automation*, 12(6):869–880, 1996.

[Brynjarsdóttir and O'Hagan 2014] Jenný Brynjarsdóttir and Anthony O'Hagan. Learning about physical parameters: The importance of model discrepancy. *Inverse Problems*, 30(11):114007, 2014.

[Choi *et al.* 2021] HeeSun Choi, Cindy Crump, Christian Duriez, Asher Elmquist, Gregory Hager, David Han, Frank Hearl, Jessica Hodgins, Abhinandan Jain, Frederick Leve, Chen Li, Franziska Meier, Dan Negrut, Ludovic Righetti, Alberto Rodriguez, Jie Tan, and Jeff Trinkle. On the use of simulation in robotics: Opportunities, challenges, and suggestions for moving forward. *Proceedings of the National Academy of Sciences*, 118(1):e1907856118, 2021.

[Collins *et al.* 2021] Jack Collins, Shelvin Chand, Anthony Vanderkop, and David Howard. A review of physics simulators for robotic applications. *IEEE Access*, 9:51416–51431, 2021.

[Erez *et al.* 2015] Tom Erez, Yuval Tassa, and Emanuel Todorov. Simulation tools for model-based robotics: Comparison of Bullet, Havok, MuJoCo, ODE and PhysX. In *IEEE International Conference on Robotics and Automation (ICRA)*, pages 4397–4404, 2015.

[Fritzson 2014] Peter Fritzson. *Principles of Object-Oriented Modeling and Simulation with Modelica 3.3: A Cyber-Physical Approach*. Wiley-IEEE Press, Hoboken, NJ, 2nd edition, 2014.

[Fuller *et al.* 2020] Aidan Fuller, Zhong Fan, Charles Day, and Chris Barlow. Digital twin: Enabling technologies, challenges and open research. *IEEE Access*, 8:108952–108971, 2020.

[Gelman and Loken 2014] Andrew Gelman and Eric Loken. The statistical crisis in science. *American Scientist*, 102(6):460–465, 2014.

[Grieves and Vickers 2017] Michael Grieves and John Vickers. Digital twin: Mitigating unpredictable, undesirable emergent behavior in complex systems. In Franz-Josef Kahlen, Shannon Flumerfelt, and Anabela Alves, editors, *Transdisciplinary Perspectives on Complex Systems*, pages 85–113. Springer, Cham, 2017.

[Ho *et al.* 1975] Chung-Wen Ho, Albert E. Ruehli, and Pierce A. Brennan. The modified nodal approach to network analysis. *IEEE Transactions on Circuits and Systems*, 22(6):504–509, 1975.

[Hughes and Drury 2019] Austin Hughes and Bill Drury. *Electric Motors and Drives: Fundamentals, Types and Applications*. Newnes, Oxford, 5th edition, 2019.

[Hwangbo *et al.* 2019] Jemin Hwangbo, Joonho Lee, Alexey Dosovitskiy, Dario Bellicoso, Vassilios Tsounis, Vladlen Koltun, and Marco Hutter. Learning agile and dynamic motor skills for legged robots. *Science Robotics*, 4(26):eaau5872, 2019.

[Jakobi *et al.* 1995] Nick Jakobi, Phil Husbands, and Inman Harvey. Noise and the reality gap: The use of simulation in evolutionary robotics. In *Advances in Artificial Life (ECAL 1995)*, volume 929 of *Lecture Notes in Computer Science*, pages 704–720. Springer, 1995.

[Jones *et al.* 2020] David Jones, Chris Snider, Aydin Nassehi, Jason Yon, and Ben Hicks. Characterising the digital twin: A systematic literature review. *CIRP Journal of Manufacturing Science and Technology*, 29:36–52, 2020.

[Kennedy and O'Hagan 2001] Marc C. Kennedy and Anthony O'Hagan. Bayesian calibration of computer models. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 63(3):425–464, 2001.

[Kitano *et al.* 1997] Hiroaki Kitano, Minoru Asada, Yasuo Kuniyoshi, Itsuki Noda, and Eiichi Osawa. RoboCup: The robot world cup initiative. In *Proceedings of the First International Conference on Autonomous Agents*, pages 340–347, 1997.

[Koenig and Howard 2004] Nathan Koenig and Andrew Howard. Design and use paradigms for Gazebo, an open-source multi-robot simulator. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 2149–2154, 2004.

[Koos *et al.* 2013] Sylvain Koos, Jean-Baptiste Mouret, and Stéphane Doncieux. The transferability approach: Crossing the reality gap in evolutionary robotics. *IEEE Transactions on Evolutionary Computation*, 17(1):122–145, 2013.

[Kritzinger *et al.* 2018] Werner Kritzinger, Matthias Karner, Georg Traar, Jan Henjes, and Wilfried Sihn. Digital twin in manufacturing: A categorical literature review and classification. *IFAC-PapersOnLine*, 51(11):1016–1022, 2018.

[Ljung 1999] Lennart Ljung. *System Identification: Theory for the User*. Prentice Hall, Upper Saddle River, NJ, 2nd edition, 1999.

[Muratore *et al.* 2022] Fabio Muratore, Fabio Ramos, Greg Turk, Wenhao Yu, Michael Gienger, and Jan Peters. Robot learning from randomized simulations: A review. *Frontiers in Robotics and AI*, 9:799893, 2022.

[Nagel and Pederson 1973] Laurence W. Nagel and Donald O. Pederson. *SPICE (Simulation Program with Integrated Circuit Emphasis)*. Technical Report UCB/ERL M382, Electronics Research Laboratory, University of California, Berkeley, 1973.

[Nosek *et al.* 2018] Brian A. Nosek, Charles R. Ebersole, Alexander C. DeHaven, and David T. Mellor. The preregistration revolution. *Proceedings of the National Academy of Sciences*, 115(11):2600–2606, 2018.

[NVIDIA AI IOT 2019] NVIDIA AI IOT. *JetBot: An Educational AI Robot Based on NVIDIA Jetson Nano*. https://github.com/NVIDIA-AI-IOT/jetbot, 2019. Accessed 1 November 2026.

[Oberkampf and Roy 2010] William L. Oberkampf and Christopher J. Roy. *Verification and Validation in Scientific Computing*. Cambridge University Press, Cambridge, 2010.

[Pacejka 2012] Hans B. Pacejka. *Tire and Vehicle Dynamics*. Butterworth-Heinemann, Oxford, 3rd edition, 2012.

[Peng *et al.* 2018] Xue Bin Peng, Marcin Andrychowicz, Wojciech Zaremba, and Pieter Abbeel. Sim-to-real transfer of robotic control with dynamics randomization. In *IEEE International Conference on Robotics and Automation (ICRA)*, pages 3803–3810, 2018.

[Rasheed *et al.* 2020] Adil Rasheed, Omer San, and Trond Kvamsdal. Digital twin: Values, challenges and enablers from a modeling perspective. *IEEE Access*, 8:21980–22012, 2020.

[Sargent 2013] Robert G. Sargent. Verification and validation of simulation models. *Journal of Simulation*, 7(1):12–24, 2013.

[Siegwart *et al.* 2011] Roland Siegwart, Illah R. Nourbakhsh, and Davide Scaramuzza. *Introduction to Autonomous Mobile Robots*. MIT Press, Cambridge, MA, 2nd edition, 2011.

[Tan *et al.* 2018] Jie Tan, Tingnan Zhang, Erwin Coumans, Atil Iscen, Yunfei Bai, Danijar Hafner, Steven Bohez, and Vincent Vanhoucke. Sim-to-real: Learning agile locomotion for quadruped robots. In *Robotics: Science and Systems (RSS)*, 2018.

[Tao *et al.* 2019] Fei Tao, He Zhang, Ang Liu, and A. Y. C. Nee. Digital twin in industry: State-of-the-art. *IEEE Transactions on Industrial Informatics*, 15(4):2405–2415, 2019.

[Tobin *et al.* 2017] Josh Tobin, Rachel Fong, Alex Ray, Jonas Schneider, Wojciech Zaremba, and Pieter Abbeel. Domain randomization for transferring deep neural networks from simulation to the real world. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 23–30, 2017.

[Todorov *et al.* 2012] Emanuel Todorov, Tom Erez, and Yuval Tassa. MuJoCo: A physics engine for model-based control. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 5026–5033, 2012.

[Waveshare 2019] Waveshare. *JetBot AI Kit*. https://www.waveshare.com/wiki/JetBot_AI_Kit, 2019. Accessed 1 November 2026.

[Zhao *et al.* 2020] Wenshuai Zhao, Jorge Peña Queralta, and Tomi Westerlund. Sim-to-real transfer in deep reinforcement learning for robotics: A survey. In *IEEE Symposium Series on Computational Intelligence (SSCI)*, pages 737–744, 2020.

Wits University Faculty of Science post-graduate student AI declaration

I understand that the use of generative AI tools (such as ChatGPT or similar) without explicitly declaring such use constitutes a form of plagiarism and is classified by Wits University as academic misconduct.

I declare that in the course of conducting the research towards my degree or in the preparation of this thesis/dissertation/research report (select one by marking with an X):

I **did not** make use of generative AI tools [ ]

I **did** make use of generative AI tools for the following (tick all that apply):

1.	Idea Generation (research problem/design, hypothesis)	
2.	Sourcing Related Work (summarising, identifying sources)	X
3.	Methods and Experiment Design (experiment setup, model tuning)	
4.	Data Analysis (presentation, coding, interpretation)	X
5.	Theoretical Development (theorem proving, conceptual analysis)	X
6.	Code Development (generating algorithms, writing scripts)	X
7.	Presentation (rendering graphics, formatting)	X
8.	Editing (grammar, readability)	X
9.	Writing (text generation, document structuring)	X
10.	Citation Formatting (structuring, organising)	X

If other uses were involved, please specify below:

Generative AI tool used (list all)	**Used for?**
Claude (Anthropic), via Claude Code	All uses ticked above (items 2 and 4 to 10)

If generative AI tools were used as an integral part of the experimental design or in the direct execution of my research, I confirm that details of this use are clearly outlined in the relevant experimental/methodology chapters of my thesis/dissertation/research report.

Student number: 2596286

Candidate signature: *Muthuphei Mukhunyeledzi*

Date: 2 November 2026